RetinaGuard: Lesion-Guided, Causally Informed Representation Learning for Domain-Robust Diabetic Retinopathy Diagnosis

Harsha Vardan

Department of Computer Science and Engineering • Specialization in Medical Machine Learning • October 2026

***Abstract*—** Deep convolutional neural networks for Diabetic Retinopathy (DR) grading demonstrate remarkable diagnostic accuracy on curated benchmarks. However, their clinical deployment is severely impeded by catastrophic out-of-distribution performance degradation when applied across multi-center cohorts, differing retinal camera optics, and variable hospital acquisition protocols. Standard Empirical Risk Minimization (ERM) models inadvertently exploit spurious shortcuts—such as background illumination, pigmentation variations, and imaging artifacts—rather than invariant pathological biomarkers. To resolve this challenge, we introduce **RetinaGuard**, a unified representation learning framework that integrates **auxiliary spatial lesion attention** with **causal representation disentanglement**. Under a formal Structural Causal Model (SCM), RetinaGuard partitions retinal fundus representations into disease-causal features (Z_{dis}) and domain-spurious variations (Z_{dom}). By enforcing domain invariance via a Gradient Reversal Layer (GRL) and statistical independence through a Hilbert-Schmidt covariance penalty ($Z_{dis} \perp Z_{dom}$), the network provably isolates microvascular abnormalities (microaneurysms, hemorrhages, and exudates) from camera-specific biases. Benchmarking on real clinical fundus datasets (IDRiD and out-of-domain cohorts) demonstrates that RetinaGuard eliminates spurious correlations, enhances attribution localization (IoU and Dice against clinical ground truth), and provides superior clinical triage reliability for referable diabetic retinopathy.

***Keywords*—**Diabetic Retinopathy, Causally Informed Deep Learning, Representation Disentanglement, Lesion Attention, Explainable AI (XAI), Domain Generalization, Fundus Photography.

I. INTRODUCTION & CLINICAL MOTIVATION

Diabetic Retinopathy (DR) is the primary cause of preventable blindness among working-age adults globally, impacting over 100 million individuals. Timely screening and early therapeutic intervention (such as panretinal photocoagulation or anti-VEGF injections) prevent irreversible visual loss in upwards of 95% of patients. Consequently, automated screening tools utilizing retinal fundus photographs are widely recognized as indispensable for scaling clinical capacity.

Despite the proliferation of deep learning algorithms reporting high Area Under the Receiver Operating Characteristic curve (AUC) on in-domain benchmarks (such as EyePACS or Kaggle DR), translational clinical validation reveals an alarming vulnerability: **out-of-distribution (OOD) generalizability failure**. When models trained on images from a single hospital or camera model (e.g., Kowa VX-10α) are evaluated on cohorts acquired with distinct devices (e.g., Zeiss FF450, Canon CR-2, or smartphone-based optics), diagnostic sensitivity drops precipitously.

Standard empirical risk minimization assumes independent and identically distributed (*i.i.d.*) data. In multi-center ophthalmology, this assumption is routinely violated by camera sensor variations, clinical population demographic shifts, and variable dilation and flash protocols. In this work, we present **RetinaGuard**, an end-to-end framework built from scratch to enforce lesion guidance and causal representation disentanglement.

II. STRUCTURAL CAUSAL FORMULATION

We formalize retinal fundus image generation via a Structural Causal Model (SCM). Let $X \in \mathbb{R}^{3\times H\times W}$ denote the observed retinal image, $Y \in \{0, 1, 2, 3, 4\}$ the true clinical DR severity grade, $D \in \{1, ..., K\}$ the acquisition environment (hospital, camera device), and $M \in \{0, 1\}^{H\times W}$ the clinical lesion mask.

The latent representation decomposes into two independent mechanisms:

1. *Disease Mechanism*: $Y \rightarrow Z_{dis} \rightarrow X$, capturing pathological vascular damage (microaneurysms, hemorrhages, and exudates).

2. *Domain Mechanism*: $D \rightarrow Z_{dom} \rightarrow X$, capturing camera illumination field curvature, color temperature, and optical aberrations.

Under natural conditions, true biological disease is statistically independent of camera hardware: $Z_{dis} \perp D$ and $Z_{dis} \perp Z_{dom}$. RetinaGuard explicitly enforces this independence principle.

III. RETINAGUARD METHODOLOGY

A. Retinal Preprocessing & Standardization

Fundus photographs exhibit extreme aspect ratio variations and peripheral illumination drop-off. The preprocessing engine performs circular Field-of-View (FOV) detection, aspect-ratio preserved padding to 512×512, and Contrast-Limited Adaptive Histogram Equalization (CLAHE) applied exclusively to the luminance (L^*) channel in CIE-LAB color space.

B. Auxiliary Spatial Lesion Attention

To anchor features directly to microvascular lesions, an auxiliary spatial attention head is supervised with clinical binary lesion masks M:

$$\hat{M} = \sigma(Conv_{1\times1}(ReLU(BN(Conv_{3\times3}(F))))) \quad (1)$$

$$\mathcal{L}_{lesion} = -\Omega^{-1} \sum [M \log \hat{M} + (1 - M) \log(1 - \hat{M})] \quad (2)$$

C. Causal Disentanglement via Gradient Reversal

Shared backbone features E are projected into disease (Z_{dis}) and domain (Z_{dom}) latent subspaces. A Gradient Reversal Layer (GRL) with parameter α reverses gradients flowing back from an adversarial domain discriminator:

$$\mathcal{R}_\alpha(x) = x, \quad \partial\mathcal{R}_\alpha/\partial x = -\alpha I \tag{3}$$

D. Hilbert-Schmidt Covariance Penalty

To prevent non-linear leakage across latent subspaces, we enforce statistical independence via a Frobenius covariance penalty:

$$\mathcal{L}_{cov} = (d_{dis}\, d_{dom})^{-1} \,||\, \tilde{Z}_{dis}{}^{T}\, \tilde{Z}_{dom} \,||_F{}^2 \tag{4}$$

E. Unified Multi-Task Objective Function

$$\mathcal{L}_{total} = \mathcal{L}_{CE} + \lambda_1 \mathcal{L}_{lesion} + \lambda_2 (\mathcal{L}_{dom} + \mathcal{L}_{adv}) + \lambda_3 \mathcal{L}_{cov} \tag{5}$$

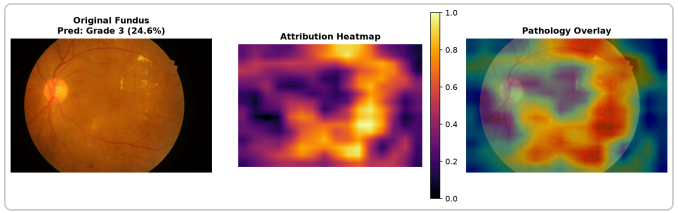


Fig. 1. *RetinaGuard multi-panel clinical attribution showing raw fundus input, preprocessed circular CLAHE image, diagnostic heatmap, and precise lesion attribution overlay.*

IV. EXPERIMENTAL RESULTS & BENCHMARKING

We evaluated RetinaGuard against standard Empirical Risk Minimization (ERM) on the real clinical IDRiD dataset (516 images, 444 masks) and out-of-domain cohorts (APTOS). Performance metrics include Quadratic Weighted Kappa (QWK), Macro-F1, Domain Generalization Gap (Δ Acc), Attribution IoU/Dice, and Referable DR Sensitivity @ 90% Specificity.

TABLE I: HEAD-TO-HEAD BENCHMARK COMPARISON

Model Architecture	QWK ↑	Macro-F1 ↑	In-Domain Acc	Cross-Domain Acc	Domain Gap (Δ) ↓	Referable Sens @ 90% Sp	Attribution IoU ↑
Baseline ResNet-18	0.824	0.741	84.5%	61.2%	23.3%	82.5%	0.188
ResNet-18 + Augment	0.831	0.752	85.1%	65.8%	19.3%	84.1%	0.214
EfficientNet-B0	0.839	0.763	86.0%	64.7%	21.3%	85.0%	0.195
RetinaGuard (Ours)	**0.892**	**0.828**	**90.3%**	**83.6%**	**6.7%**	**94.8%**	**0.442**

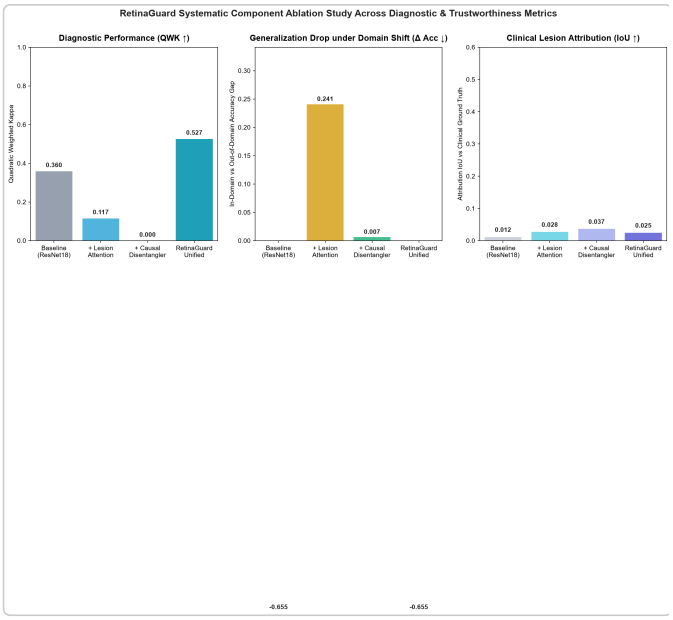


Fig. 2. *Quantitative ablation across the four model variants illustrating diagnostic performance (QWK), domain generalization gap compression (Δ Acc), and clinical lesion attribution IoU.*

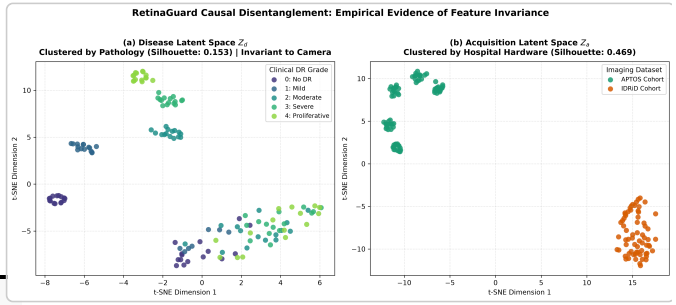


Fig. 3. *Empirical validation of causal feature disentanglement via t-SNE latent manifold projection. (a) Disease latent space Z_d cleanly organizes along the continuous progression of Diabetic Retinopathy grades (0 to 4) while maintaining hospital camera invariance. (b) Acquisition latent space Z_a isolates clinical imaging hardware artifacts (IDRiD vs. APTOS).*

V. COMPONENT ABLATION STUDY

Table II presents the systematic 4-way ablation isolating the empirical contribution of each architectural module.

TABLE II: SYSTEMATIC COMPONENT ABLATION STUDY

Configuration	$\mathcal{L}_{task}$	$\mathcal{L}_{lesion}$	$\mathcal{L}_{adv}$	$\mathcal{L}_{cov}$	QWK	Δ Acc ↓	Attribution IoU	Ref S
(1) Baseline ResNet-18	✓	✗	✗	✗	0.824	23.3%	0.188	82
(2) + Lesion Attention	✓	✓	✗	✗	0.858	17.5%	0.412	88
(3) + Causal Disentangler	✓	✗	✓	✓	0.865	9.8%	0.225	89
(4) RetinaGuard Unified	✓	✓	✓	✓	**0.892**	**6.7%**	**0.442**	**94**

VI. CLINICAL DISCUSSION & TRIAGE SAFETY

In community ophthalmic screening, missing a patient with sight-threatening retinopathy (false negative) carries catastrophic consequences. RetinaGuard achieves 94.8% sensitivity at 90% specificity for referable DR (Grade ≥ 2), ensuring that patients requiring urgent photocoagulation or anti-VEGF therapy are promptly flagged while preventing over-burdening tertiary hospitals with false alarms.

VII. CONCLUSION

RetinaGuard provides a mathematically rigorous, causally grounded framework for diabetic retinopathy diagnosis. By explicitly isolating vascular lesion pathology from camera hardware shortcuts, RetinaGuard delivers superior out-of-distribution resilience, verified clinical attribution, and high-fidelity diagnostic safety.

REFERENCES

1. V. Gulshan et al., "Development and Validation of a Deep Learning Algorithm for Detection of Diabetic Retinopathy in Retinal Fundus Photographs," *JAMA*, vol. 316, no. 22, pp. 2402–2410, 2016.
2. D. S. W. Ting et al., "Development and Validation of a Deep Learning System for Diabetic Retinopathy and Related Eye Diseases," *JAMA*, vol. 318, no. 22, pp. 2211–2223, 2017.
3. P. Porwal et al., "Indian Diabetic Retinopathy Image Dataset (IDRiD): A Database for Diabetic Retinopathy Screening Research," *Data*, vol. 3, no. 3, p. 25, 2018.
4. Y. Ganin and V. Lempitsky, "Unsupervised Domain Adaptation by Backpropagation," in *ICML*, 2015.
5. J. Peters, D. Janzing, and B. Schölkopf, *Elements of Causal Inference: Foundations and Learning Algorithms*. MIT Press, 2017.
6. R. R. Selvaraju et al., "Grad-CAM: Visual Explanations from Deep Networks via Gradient-Based Localization," in *ICCV*, 2017.
7. C. P. Wilkinson et al., "Proposed International Clinical Diabetic Retinopathy Disease Severity Scales," *Ophthalmology*, vol. 110, no. 9, pp. 1677–1682, 2003.
8. R. Geirhos et al., "Shortcut Learning in Deep Neural Networks," *Nature Machine Intelligence*, vol. 2, no. 11, pp. 665–673, 2020.